

Decoder report — Dashcam Video Integrity

The decoder verifies a video against the hashes the encoder stored in Supabase. It answers: does it match, which trip, where in the trip, what percentage matches, which sections were modified, is it a bit-exact original segment, and how long the check took.

1. Install and run

Requirements: Python ≥ 3.11 , and `ffmpeg` only for the evaluation scripts.

macOS / Linux

```
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
cp decoder/.env.example decoder/.env      # SUPABASE_URL and SUPABASE_ANON_KEY
uvicorn decoder.main:app --reload          # open http://localhost:8000
python -m pytest                           # tests
```

Windows (PowerShell)

```
python -m venv .venv
.venv\Scripts\activate
pip install -r requirements.txt
copy decoder\.env.example decoder\.env    # edit it with Notepad
uvicorn decoder.main:app --reload
python -m pytest
```

Using it: choose a trip (or "Auto — search all trips"), choose a video, press Verify. The page shows the verdict, the result table and a bar with one box per 0.5 s (green = matched, red = changed, grey = no stored sample at that time). "Purge" deletes server data older than N hours.

API: POST `/api/verify` (form: video, optional trip_id), GET `/api/trips`, POST `/api/purge`.

Evaluation (needs `ffmpeg`; `brew install ffmpeg` / `winget install ffmpeg`):

```
python eval/make_variants.py eval/data/tripA.mp4 eval/data/tripB.mp4
python -m eval.metric_study      # metric comparison, writes decoder/thresholds.json
python -m eval.video_eval <trip id>
```

2. How it works (and why)

1. **Exact check:** SHA-256 of the uploaded file is looked up in `segments`. A hit means the file is byte-for-byte a recorded segment ("bit-exact original segment N"). Any change breaks it.
2. **Sampling:** the video is read frame by frame with OpenCV and one pHash is taken every 500 ms of video time (same rate as the encoder). Frames are read sequentially, not by seeking, because browser WebM files have no index.
3. **Sliding window:** candidate offsets put the first video sample on every stored sample (video starts inside the trip) and the first stored sample on every video sample (extra content before the trip). For each offset, each video sample is compared with the stored sample **closest in time** (`t_ms`, up to 600 ms away). The offset with the lowest mean Hamming distance is the **position in the trip**. Missing samples count as 32 bits, the distance between unrelated hashes.

4. **Per sample:** matched if distance $\leq$ threshold (16 bits, from the evaluation below). **Two or more** unmatched samples in a row = a **modified section** (from-to seconds).
5. **Verdict:** **AUTHENTIC** (≥ 90 % matched and no modified section), **MODIFIED** (30–90 % or any modified section), **NO MATCH** (< 30 %).
6. Without a trip ID the video is compared with every trip and the best one wins. A bit-exact SHA-256 match always selects its own trip.

Why pHash in the decoder although cosine scored slightly higher (below): a pHash is 64 bits (16 hex chars) per sample and is compared with a fast Hamming distance. The cosine metric needs a 256-value vector per sample. pHash is also the same algorithm in the browser and in Python (verified by a test). Its accuracy is almost identical (98.9 % vs 99.8 %) with **0 false positives**.

3. Results

Data: trip A = tripA.mp4 (84 s dashcam video, recorded with the encoder in file mode, trip a2280b49), trip B = tripB.mp4 (10.4 s, trip e2b53bdc). Variants were made from a 720p copy of trip A (4K processing was too slow; the hashes work on 32x32 images).

3.1 Metric comparison (eval/results/thresholds.csv)

1008 positive pairs (frame of A vs the same timestamp in a content-preserving variant, 1 per second, 12 variants) and 84 negative pairs (frame of A vs frame of B). Threshold = best balanced accuracy.

Metric	Threshold	Accuracy	True pos. rate	False pos. rate	Mean same content	Mean other trip
cosine (16x16 gray)	0.048	0.998	0.996	0.000	0.005	0.117
pHash Hamming	16	0.989	0.978	0.000	2.38	28.19
pHash normalized	0.250	0.989	0.978	0.000	0.037	0.440
wHash Hamming	10	0.989	0.990	0.012	1.09	20.45
wHash normalized	0.156	0.989	0.990	0.012	0.017	0.320
dHash Hamming	19	0.985	0.969	0.000	2.48	30.69
dHash normalized	0.297	0.985	0.969	0.000	0.039	0.480
aHash Hamming	7	0.979	0.982	0.024	0.92	12.87
aHash normalized	0.109	0.979	0.982	0.024	0.014	0.201
L1 (16x16 gray)	0.120	0.958	0.917	0.000	0.034	0.195
L2 (16x16 gray)	2.64	0.957	0.915	0.000	0.668	3.959
TLSH (JPEG bytes)	214	0.885	0.829	0.060	194.9	240.8
ssdeep (JPEG bytes)	—	0.500	—	—	100	100

Normalized Hamming = Hamming / 64, so it gives the same ranking. **ssdeep and TLSH do not work** on video frames: they hash the compressed bytes, and any re-encoding changes all the bytes (ssdeep similarity was 0 for every pair). All histograms: eval/results/hist_*.png.




Figure 1 — Distance histograms: same content (blue) vs other trip (orange), dashed line = chosen threshold. pHash, dHash and cosine separate the two groups; TLSh overlaps.

Mean Hamming distance (bits) per variant (`eval/results/mean_distance_per_variant.csv`):

Variant	pHash	dHash	aHash	wHash
original (re-encoded)	0.12	0.12	0.06	0.05
H.265	0.10	0.37	0.05	0.17
640x360	0.17	0.37	0.05	0.12
Gaussian noise	0.31	0.38	0.11	0.17
salt & pepper 5 %	0.38	0.90	0.19	0.29
CRF 40	0.79	1.48	0.23	0.45
300 kbit/s	0.81	1.80	0.27	0.57
10 fps	1.36	0.88	0.23	0.52
brightness +0.2	2.71	1.23	1.58	0.64
contrast x1.5	2.98	3.24	1.18	1.88
logo overlay	4.05	0.68	0.74	0.26
crop 80 %	14.76	18.30	6.36	8.00
trip A vs trip B	28.19	30.69	12.87	20.45

3.2 Decoder on whole videos (`eval/results/video_eval.md`)

The real decoder (hashes of the variants in Python vs the hashes the **browser** stored for trip A), pHash threshold 16 bits.

Variant	Expected	Verdict	Matched %	Mean dist.	Position (s)	Modified sections	Time (ms)
original	AUTHENTIC	AUTHENTIC	100	1.25	0.1	—	3261

Variant	Expected	Verdict	Matched %	Mean dist.	Position (s)	Modified sections	Time (ms)
trim start (–10 s)	AUTHENTIC	AUTHENTIC	100	1.23	10.1	–	2955
trim end (–10 s)	AUTHENTIC	AUTHENTIC	100	1.24	0.1	–	3225
extract middle (20 s)	AUTHENTIC	AUTHENTIC	100	1.10	32.1	–	833
shift (3 s black at start)	MODIFIED	MODIFIED	96.6	2.36	–2.9	0.0–3.0 s	3646
speed ×1.05	MODIFIED	MODIFIED	66.2	13.28	3.1	10 sections	3332
10 fps	AUTHENTIC	AUTHENTIC	100	1.32	0.1	–	2364
delete 1 s	MODIFIED	MODIFIED	88.6	6.39	1.1	4.0–11.0; 14.0–15.0; 40.5–41.5 s	3528
duplicate 1 s	MODIFIED	MODIFIED	87.6	6.45	–0.9	5.0–12.0; 13.0–14.0; 41.5–43.0 s	3515
H.265	AUTHENTIC	AUTHENTIC	100	1.38	0.1	–	3382
300 kbit/s	AUTHENTIC	AUTHENTIC	100	1.61	0.1	–	2784
640×360	AUTHENTIC	AUTHENTIC	100	1.31	0.1	–	901
CRF 40	AUTHENTIC	AUTHENTIC	100	1.56	0.1	–	2765
brightness +0.2	AUTHENTIC	AUTHENTIC	100	3.08	0.1	–	3462
contrast ×1.5	AUTHENTIC	AUTHENTIC	100	3.57	0.1	–	3587
salt & pepper 5 %	AUTHENTIC	AUTHENTIC	100	1.36	0.1	–	7604
Gaussian noise	AUTHENTIC	AUTHENTIC	100	1.39	0.1	–	6856
logo overlay	AUTHENTIC	AUTHENTIC	100	4.13	0.1	–	3539
crop 80 %	AUTHENTIC	MODIFIED x	71.4	14.69	0.1	11 sections	1924
other trip (B)	NO MATCH	NO MATCH	19.0	23.52	7.6	–	500
partial replacement (5 s of B)	MODIFIED	MODIFIED	93.5	3.58	0.1	42.0–47.5 s	3368




Figure 2 — Matched samples per variant (pHash threshold 16 bits). Dashed lines: verdict limits.

Summary: 21 variants, **20 correct verdicts**. True matches 20, **missed matches 0, false matches 0**, modifications detected **5 / 5**. Processing time 0.5–7.6 s per variant (10–87 s of 720p video, laptop CPU). Position checks: trimming 10 s gives 10.1 s; the middle extract (cut at 32.0 s) gives 32.1 s; the 3 s black intro gives –2.9 s. The replaced 5 s (42.0–47.0 s) are reported as 42.0–47.5 s.

Real browser segment (manual check): a WebM segment downloaded from the encoder (`trip-72d0f447-seg0.webm`) → exact SHA-256 match with segment 0, **AUTHENTIC**, position 0.0 s, mean distance 3.0 bits. The JS and Python hashes of the same frames differ by 0–2 bits on most samples.

4. Limitations

- **Crop** changes the whole frame for a global hash → crop 80 % is not recognised as authentic (14.8 bits mean pHash distance, close to the 16-bit threshold). This is the only wrong verdict.
- **Only one offset per video:** after a cut or a duplicated second, one half of the video is aligned and the other is 1 s off. The video is correctly flagged **MODIFIED**, but some of the reported sections are in fast-moving parts of the misaligned half, not exactly at the edit (e.g. delete 1 s at 42 s → sections 4–11 s and 40.5–41.5 s). The same happens with **speed changes**.
- **Visually similar trips** (e.g. two webcam recordings of the same desk) look the same to perceptual hashes; the decoder may confuse them. Real dashcam trips differ a lot (28 bits apart on average).
- "Other trip" got 19 % matched samples, under the 30 % **NO MATCH** limit but not by much (B is only 10 s).
- Big logos or overlays covering a large part of the frame will raise the distance (a small logo gave 4.1 bits).
- Thresholds come from one pair of videos; more varied footage (night, rain) would give more reliable values.
- No authentication on the API (prototype).